

Modeling - Seoul Bike Sharing Data

Istruzioni: lanciare tutti i code chunks con Ctrl+Alt+R in modo da evitare di lanciare i codici che stimano i modelli (i modelli sono già stati salvati nella cartella models)

Carico pacchetti:

Carico dati:

Modelli “all features”

Random forest 1

Random forest stimata per feature selection iniziale

Codice per stima modello (non necessario eseguire)

Save/Load model:

```
bike_rf1 <- readRDS("models/bike_rf1.rda")
```

```
paste("Validation RMSE: ", RMSE(round(predict(bike_rf1, bike_train)$predictions), bike_train$rented_bike_count))
```

RMSE

```
## [1] "Validation RMSE: 86.7527482065492"
```

```
paste("Validation RMSE: ", RMSE(round(predict(bike_rf1, bike_valid)$predictions), bike_valid$rented_bike_count))
```

```
## [1] "Validation RMSE: 194.680474112414"
```

```
paste("Testing RMSE: ", RMSE(round(predict(bike_rf1, bike_test)$predictions), bike_test$rented_bike_count))
```

```
## [1] "Testing RMSE: 198.944113443883"
```

Regressione lineare

Stima modello:

```
lm1 <- lm(rented_bike_count ~ ., bike_train_dummy)
```

Riassunto modello:

```
summary(lm1)
```

```
##
## Call:
## lm(formula = rented_bike_count ~ ., data = bike_train_dummy)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -1283.86  -264.38   -49.44   203.05  1969.80
##
## Coefficients: (4 not defined because of singularities)
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   -315.07194   117.92716   -2.672  0.007565 **
## hour           26.67204     0.83715   31.861 < 2e-16 ***
## temperature    19.32989     4.21261    4.589 4.55e-06 ***
## humidity       -10.25038     1.17485   -8.725 < 2e-16 ***
## wind_speed     21.18798     5.74716    3.687 0.000229 ***
## visibility      0.05535     0.01300    4.258 2.10e-05 ***
## dew_point_temperature 12.14990     4.42625    2.745 0.006069 **
## solar_radiation -84.84009     8.57337   -9.896 < 2e-16 ***
## rainfall       -53.06086     4.53112  -11.710 < 2e-16 ***
## snowfall       41.16752    12.52006    3.288 0.001014 **
## seasons_Spring -65.30545    31.25898   -2.089 0.036732 *
## seasons_Summer 196.00481    27.02258    7.253 4.56e-13 ***
## seasons_Winter -217.05726    46.03734   -4.715 2.47e-06 ***
## holiday_No.Holiday 151.05293    24.06710    6.276 3.70e-10 ***
## functioning_day_Yes 965.91655    30.58170   31.585 < 2e-16 ***
## weekday_Mon     -63.96823    19.35897   -3.304 0.000957 ***
## weekday_Sat     -77.65615    19.27358   -4.029 5.66e-05 ***
## weekday_Sun    -137.76117    19.35961   -7.116 1.24e-12 ***
## weekday_Thu     -35.18964    19.30786   -1.823 0.068418 .
## weekday_Tue     -41.57479    19.29923   -2.154 0.031261 *
## weekday_Wed     -18.75070    19.36517   -0.968 0.332947
## month_Aug       -554.20969    28.43136  -19.493 < 2e-16 ***
## month_Dec        61.93244    25.60520    2.419 0.015602 *
## month_Feb       -41.77876    26.24409   -1.592 0.111451
## month_Jan        NA         NA         NA         NA
## month_Jul       -401.85129    27.19479  -14.777 < 2e-16 ***
## month_Jun        NA         NA         NA         NA
## month_Mar       -73.45922    26.16311   -2.808 0.005004 **
## month_May       138.00549    26.82785    5.144 2.77e-07 ***
## month_Nov       78.94940    34.80561    2.268 0.023345 *
## month_Oct       170.50512    28.29167    6.027 1.77e-09 ***
## month_Sep        NA         NA         NA         NA
## weekend_Yes      NA         NA         NA         NA
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 410.5 on 6278 degrees of freedom
## Multiple R-squared:  0.5947, Adjusted R-squared:  0.5929
## F-statistic: 329 on 28 and 6278 DF, p-value: < 2.2e-16
```

Investighiamo le variabili per cui i coefficienti stimati della regressione lineare sono NA

Random forest dummy default

Modello applicato a dati con variabili dummy.

(Modello di default senza hyperparameter tuning)

Save/Load model:

```
bike_dummy_rf <- readRDS("models/bike_dummy_rf.rda")
```

```
default_train_rmse <- RMSE(round(predict(bike_dummy_rf, bike_train_dummy)$predictions),
                           bike_train_dummy$rented_bike_count)
default_valid_rmse <-
  RMSE(round(predict(bike_dummy_rf,
                   bike_valid_dummy)$predictions),
        bike_valid_dummy$rented_bike_count)
default_test_rmse <-
  RMSE(round(predict(bike_dummy_rf, bike_test_dummy)$predictions),
        bike_test_dummy$rented_bike_count)

paste("Training RMSE:", default_train_rmse)
```

RMSE

```
## [1] "Training RMSE: 85.0895861573404"
```

```
paste("Validation RMSE: ", default_valid_rmse)
```

```
## [1] "Validation RMSE: 193.3215407959"
```

```
paste("Testing RMSE: ", default_test_rmse)
```

```
## [1] "Testing RMSE: 195.569219548228"
```

Regressione lineare 2

Regressione lineare stimata escludendo le variabili che danno coefficienti stimati NA

Creiamo nuovo dataframe per training e test set:

```
bike_train_dummy2 <- bike_train_dummy %>%
  select(-names(which(is.na(lm1$coefficients))))

bike_valid_dummy2 <- bike_valid_dummy %>%
  select(-names(which(is.na(lm1$coefficients))))

bike_test_dummy2 <- bike_test_dummy %>%
  select(-names(which(is.na(lm1$coefficients))))
```

Stima regressione lineare:

```
lm2 <- lm(rented_bike_count ~ ., bike_train_dummy2)
```

Riassunto modello:

```
summary(lm2)
```

```
##
## Call:
## lm(formula = rented_bike_count ~ ., data = bike_train_dummy2)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -1283.86  -264.38   -49.44   203.05  1969.80
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   -315.07194   117.92716   -2.672  0.007565 **
## hour           26.67204     0.83715   31.861 < 2e-16 ***
## temperature    19.32989     4.21261    4.589 4.55e-06 ***
## humidity       -10.25038     1.17485   -8.725 < 2e-16 ***
## wind_speed     21.18798     5.74716    3.687 0.000229 ***
## visibility      0.05535     0.01300    4.258 2.10e-05 ***
## dew_point_temperature 12.14990     4.42625    2.745 0.006069 **
## solar_radiation -84.84009     8.57337   -9.896 < 2e-16 ***
## rainfall       -53.06086     4.53112  -11.710 < 2e-16 ***
## snowfall       41.16752    12.52006    3.288 0.001014 **
## seasons_Spring -65.30545    31.25898   -2.089 0.036732 *
## seasons_Summer 196.00481    27.02258    7.253 4.56e-13 ***
## seasons_Winter -217.05726    46.03734   -4.715 2.47e-06 ***
## holiday_No.Holiday 151.05293    24.06710    6.276 3.70e-10 ***
## functioning_day_Yes 965.91655    30.58170   31.585 < 2e-16 ***
## weekday_Mon     -63.96823    19.35897   -3.304 0.000957 ***
## weekday_Sat     -77.65615    19.27358   -4.029 5.66e-05 ***
## weekday_Sun    -137.76117    19.35961   -7.116 1.24e-12 ***
## weekday_Thu     -35.18964    19.30786   -1.823 0.068418 .
## weekday_Tue     -41.57479    19.29923   -2.154 0.031261 *
## weekday_Wed     -18.75070    19.36517   -0.968 0.332947
## month_Aug       -554.20969    28.43136  -19.493 < 2e-16 ***
## month_Dec        61.93244    25.60520    2.419 0.015602 *
## month_Feb       -41.77876    26.24409   -1.592 0.111451
## month_Jul       -401.85129    27.19479  -14.777 < 2e-16 ***
## month_Mar       -73.45922    26.16311   -2.808 0.005004 **
## month_May       138.00549    26.82785    5.144 2.77e-07 ***
## month_Nov        78.94940    34.80561    2.268 0.023345 *
## month_Oct       170.50512    28.29167    6.027 1.77e-09 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 410.5 on 6278 degrees of freedom
## Multiple R-squared:  0.5947, Adjusted R-squared:  0.5929
## F-statistic: 329 on 28 and 6278 DF, p-value: < 2.2e-16
```

```
paste("Training RMSE: ", RMSE(round(predict(lm2, bike_train_dummy2)), bike_train_dummy2$rented_bike_count))
```

RMSE

```
## [1] "Training RMSE: 409.525471080192"
```

```
paste("Validation RMSE: ", RMSE(round(predict(lm2, bike_valid_dummy2)), bike_valid_dummy2$rented_bike_count))
```

```
## [1] "Validation RMSE: 414.073145796431"
```

```
paste("Testing RMSE: ", RMSE(round(predict(lm2, bike_test_dummy2)), bike_test_dummy2$rented_bike_count))
```

```
## [1] "Testing RMSE: 416.657060619862"
```

Random forest 4

```
n_features <- length(bike_train_dummy) - 1

hyper_grid <- expand.grid(
  num.trees = c(100, n_features * 10, 500),
  mtry = floor(n_features * c(.05, .15, .25, .333, .4)),
  min.node.size = c(1, 3, 5, 10),
  replace = c(TRUE, FALSE),
  sample.fraction = c(.5, .63, .8),
  train_rmse = NA,
  valid_rmse = NA
)

# execute full cartesian grid search
for(i in seq_len(nrow(hyper_grid))) {
  # fit model for ith hyperparameter combination
  fit <- ranger(
    formula      = rented_bike_count ~ .,
    data         = bike_train_dummy,
    num.trees    = hyper_grid$num.trees[i],
    mtry         = hyper_grid$mtry[i],
    min.node.size = hyper_grid$min.node.size[i],
    replace      = hyper_grid$replace[i],
    sample.fraction = hyper_grid$sample.fraction[i],
    verbose      = FALSE,
    respect.unordered.factors = 'order'
  )
  # export OOB error
  hyper_grid$train_rmse[i] <- sqrt(fit$prediction.error)

  pred <- round(predict(fit, bike_valid_dummy)$predictions)

  hyper_grid$valid_rmse[i] <- RMSE(pred, bike_valid_dummy$rented_bike_count)
}
```

Hyperparameter tuning Carica hyperparameter grid da csv:

```
hyper_grid <- read_csv("models/rf_hyper_grid.csv")[-1]

## Rows: 360 Columns: 8
## -- Column specification -----
## Delimiter: ","
## dbl (7): ...1, num.trees, mtry, min.node.size, sample.fraction, train_rmse, ...
## lgl (1): replace
##
## i Use 'spec()' to retrieve the full column specification for this data.
## i Specify the column types or set 'show_col_types = FALSE' to quiet this message.

# assess top 10 models
hyper_grid %>%
  arrange(valid_rmse) %>%
  mutate(
    train_perc_gain = (default_train_rmse - train_rmse) /
      default_train_rmse * 100,
    valid_perc_gain = (default_valid_rmse - valid_rmse) /
      default_valid_rmse * 100) %>%
  head(10)

## # A tibble: 10 x 9
##   num.trees mtry min.node.size replace sample.fraction train_rmse valid_rmse
##   <dbl> <dbl>      <dbl> <lgl>      <dbl>      <dbl>      <dbl>
## 1      500   12          1 FALSE        0.8      177.      182.
## 2      100   12          1 FALSE        0.8      182.      182.
## 3      320   12          1 FALSE        0.8      177.      182.
## 4      100   12          3 FALSE        0.8      181.      183.
## 5      500   12          3 FALSE        0.8      177.      183.
## 6      320   12          3 FALSE        0.8      179.      184.
## 7      320   12          5 FALSE        0.8      180.      184.
## 8      100   12          5 FALSE        0.8      184.      184.
## 9      500   12          5 FALSE        0.8      180.      186.
## 10     320   10          1 FALSE        0.8      183.      186.
## # ... with 2 more variables: train_perc_gain <dbl>, valid_perc_gain <dbl>
```

Fit Fittiamo modello con gli iperparametri del modello migliore:

```
bike_rf4 <- ranger(
  formula      = rented_bike_count ~ .,
  data         = bike_train_dummy,
  num.trees    = 500,
  mtry         = 12,
  min.node.size = 1,
  replace      = FALSE,
  sample.fraction = 0.80,
  verbose      = FALSE,
  respect.unordered.factors = 'order',
)
```

Carico modello già stimato:

```
bike_rf4 <- readRDS("models/bike_rf4.rda")
```

```
paste("Training RMSE: ", RMSE(round(predict(bike_rf4, bike_train_dummy)$predictions), bike_train_dummy$
```

RMSE

```
## [1] "Training RMSE: 35.9265784946732"
```

```
paste("Validation RMSE: ", RMSE(round(predict(bike_rf4, bike_valid_dummy)$predictions), bike_valid_dummy$
```

```
## [1] "Validation RMSE: 182.408930031486"
```

```
paste("Testing RMSE: ", RMSE(round(predict(bike_rf4, bike_test_dummy)$predictions), bike_test_dummy$ren
```

```
## [1] "Testing RMSE: 186.777667991315"
```

Multilayer Perceptron

Vedi file “bike_MLP_nb.ipynb”

Carichiamo modello stimato:

Summary del modello:

```
summary(bike_mlp3)
```

```
## Model: "sequential"
## -----
## Layer (type)                Output Shape                Param #
## =====
## dense (Dense)                (None, 256)                 8448
## dense_1 (Dense)              (None, 128)                 32896
## dense_2 (Dense)              (None, 32)                  4128
## dense_3 (Dense)              (None, 1)                   33
## =====
## Total params: 45,505
## Trainable params: 45,505
## Non-trainable params: 0
## -----
```

Performance su validation set: RMSE = 171.2723

Modelli “selected features”

Regressione Lineare 3

Stimiamo ora un modello di regressione considerando solo alcune delle variabili più importanti individuate attraverso l'importance plot (8 variabili):

```
# no dew_point_temperature per correlazione con temperature
lm3 <- lm(rented_bike_count ~ hour + temperature + humidity + functioning_day_Yes + seasons_Winter + so
```

```
paste("Training RMSE: ", RMSE(round(predict(lm3, bike_train_dummy2)), bike_train_dummy2$rented_bike_cou
```

RMSE

```
## [1] "Training RMSE: 435.765900375691"
```

```
paste("Validation RMSE: ", RMSE(round(predict(lm3, bike_valid_dummy2)), bike_valid_dummy2$rented_bike_c
```

```
## [1] "Validation RMSE: 437.195561256628"
```

```
paste("Testing RMSE: ", RMSE(round(predict(lm3, bike_test_dummy2)), bike_test_dummy2$rented_bike_count)
```

```
## [1] "Testing RMSE: 445.996639318174"
```

Random forest 5

```
n_features <- 8
```

```
hyper_grid <- expand.grid(
  num.trees = c(100, n_features * 10, 500),
  mtry = floor(n_features * c(.05, .15, .25, .333, .4)),
  min.node.size = c(1, 3, 5, 10),
  replace = c(TRUE, FALSE),
  sample.fraction = c(.5, .63, .8),
  train_rmse = NA,
  valid_rmse = NA
)
```

```
# execute full cartesian grid search
for(i in seq_len(nrow(hyper_grid))) {
  # fit model for ith hyperparameter combination
  fit <- ranger(
    formula = rented_bike_count ~ hour+temperature+humidity+functioning_day_Yes+seasons_Winter+
      dew_point_temperature+solar_radiation+rainfall,
    data = bike_train_dummy,
    num.trees = hyper_grid$num.trees[i],
```



```

    mtry          = hyper_grid$mtry[i],
    min.node.size = hyper_grid$min.node.size[i],
    replace       = hyper_grid$replace[i],
    sample.fraction = hyper_grid$sample.fraction[i],
    verbose       = FALSE,
    respect.unordered.factors = 'order'
  )
  # export OOB error
  hyper_grid$train_rmse[i] <- sqrt(fit$prediction.error)

  pred <- round(predict(fit, bike_valid_dummy)$predictions)

  hyper_grid$valid_rmse[i] <- RMSE(pred, bike_valid_dummy$rented_bike_count)
}

```

Hyperparameter tuning Carica hyperparameter grid da csv:

```
hyper_grid2 <- read_csv("models/rf_hyper_grid2.csv")
```

```

## Rows: 360 Columns: 7
## -- Column specification -----
## Delimiter: ","
## dbl (6): num.trees, mtry, min.node.size, sample.fraction, train_rmse, valid_...
## lgl (1): replace
##
## i Use 'spec()' to retrieve the full column specification for this data.
## i Specify the column types or set 'show_col_types = FALSE' to quiet this message.

```

```

# assess top 10 models
hyper_grid2 %>%
  arrange(valid_rmse) %>%
  mutate(
    train_perc_gain = (default_train_rmse - train_rmse) /
      default_train_rmse * 100,
    valid_perc_gain = (default_valid_rmse - valid_rmse) /
      default_valid_rmse * 100) %>%
  head(10)

```

```

## # A tibble: 10 x 9
##   num.trees mtry min.node.size replace sample.fraction train_rmse valid_rmse
##   <dbl> <dbl>      <dbl> <lgl>      <dbl>      <dbl>      <dbl>
## 1     500     3           1 FALSE      0.63      241.      245.
## 2     500     3           1 TRUE       0.8       241.      245.
## 3     500     3           5 FALSE      0.8       241.      245.
## 4     100     3          10 FALSE      0.63      242.      245.
## 5      80     3           5 TRUE       0.8       245.      245.
## 6     500     3           3 FALSE      0.5       241.      245.
## 7     500     3          10 FALSE      0.8       241.      245.
## 8     500     3           3 FALSE      0.8       242.      245.
## 9     100     3           3 FALSE      0.8       245.      245.
## 10    500     3           1 FALSE      0.5       241.      245.
## # ... with 2 more variables: train_perc_gain <dbl>, valid_perc_gain <dbl>

```

Stimo il modello migliore:

```
bike_rf5 <- ranger(  
  formula = rented_bike_count ~ hour+temperature+humidity+functioning_day_Yes+seasons_Winter+  
    dew_point_temperature+solar_radiation+rainfall,  
  data      = bike_train_dummy,  
  num.trees = 500,  
  mtry      = 3,  
  min.node.size = 1,  
  replace    = FALSE,  
  sample.fraction = 0.63,  
  verbose    = FALSE,  
  respect.unordered.factors = 'order'  
)
```

Save/Load model:

```
# saveRDS(bike_rf5, "models/bike_rf5.rda")
```

```
bike_rf5 <- readRDS("models/bike_rf5.rda")
```

```
paste("Training RMSE: ", RMSE(round(predict(bike_rf5, bike_train_dummy)$predictions), bike_train_dummy$rented_bike_count))
```

RMSE

```
## [1] "Training RMSE: 115.358932760696"
```

```
paste("Validation RMSE: ", RMSE(round(predict(bike_rf5, bike_valid_dummy)$predictions), bike_valid_dummy$rented_bike_count))
```

```
## [1] "Validation RMSE: 244.518497250782"
```

```
paste("Testing RMSE: ", RMSE(round(predict(bike_rf5, bike_test_dummy)$predictions), bike_test_dummy$rented_bike_count))
```

```
## [1] "Testing RMSE: 251.574078798448"
```

MLP

Summary del modello:

```
summary(bike_mlp3)
```

```
## Model: "sequential"
```

```
## -----  
## Layer (type)                Output Shape          Param #  
## =====  
## dense (Dense)                (None, 256)           8448  
## dense_1 (Dense)              (None, 128)          32896
```

```

## dense_2 (Dense)                (None, 32)                4128
## dense_3 (Dense)                (None, 1)                 33
## =====
## Total params: 45,505
## Trainable params: 45,505
## Non-trainable params: 0
## -----

```

MLP

Summary del modello:

```
summary(bike_mlp4)
```

```

## Model: "sequential"
## -----
## Layer (type)                Output Shape                Param #
## =====
## dense (Dense)                (None, 64)                  576
## dense_1 (Dense)              (None, 32)                  2080
## dense_2 (Dense)              (None, 8)                   264
## dense_3 (Dense)              (None, 1)                   9
## =====
## Total params: 2,929
## Trainable params: 2,929
## Non-trainable params: 0
## -----

```